

```
In [2]: # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

/usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning: pandas.util.testing is deprecated. Use the functions in the public API at pandas.testing instead.
import pandas.util.testing as tm

```
In [0]: %matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()

    fig.canvas.draw()
```

```
In [4]: # the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 0s 0us/step

```
In [5]: print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

```
In [0]: # if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
In [7]: # after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

```
In [8]: # An example data point
print(X_train[0])
```

[illegible]

```
In [0]: # if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
X_test = X_test/255
```

```
In [10]: # example data point after normlizing  
print(X_train[0])
```

[illegible]

[illegible]

```
In [11]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])

# Lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])

Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

Softmax classifier

```
In [0]: # https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the Layer, and
# bias is a bias vector created by the Layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation
```

```
In [0]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

```
In [43]: X_train.shape[1]
```

```
Out[43]: 784
```

1. MLP 2-Hidden layer architecture

1.1. MLP + ReLU + ADAM

```
In [29]: model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history1 = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 512)	401920
dense_6 (Dense)	(None, 256)	131328
dense_7 (Dense)	(None, 10)	2570

Total params: 535,818
Trainable params: 535,818
Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 109us/step - loss: 0.2122 - accuracy: 0.9362 - val_loss: 0.1199
- val_accuracy: 0.9620

Epoch 2/20

60000/60000 [=====] - 7s 109us/step - loss: 0.0793 - accuracy: 0.9759 - val_loss: 0.0749
- val_accuracy: 0.9760

Epoch 3/20

60000/60000 [=====] - 6s 103us/step - loss: 0.0512 - accuracy: 0.9841 - val_loss: 0.0709
- val_accuracy: 0.9777

Epoch 4/20

60000/60000 [=====] - 6s 105us/step - loss: 0.0336 - accuracy: 0.9893 - val_loss: 0.0733
- val_accuracy: 0.9777

Epoch 5/20

60000/60000 [=====] - 6s 105us/step - loss: 0.0267 - accuracy: 0.9914 - val_loss: 0.0747
- val_accuracy: 0.9777

Epoch 6/20

60000/60000 [=====] - 6s 106us/step - loss: 0.0175 - accuracy: 0.9947 - val_loss: 0.0910
- val_accuracy: 0.9754

Epoch 7/20

60000/60000 [=====] - 6s 104us/step - loss: 0.0170 - accuracy: 0.9942 - val_loss: 0.0713
- val_accuracy: 0.9796

Epoch 8/20

60000/60000 [=====] - 6s 105us/step - loss: 0.0145 - accuracy: 0.9953 - val_loss: 0.0985
- val_accuracy: 0.9738

Epoch 9/20

60000/60000 [=====] - 6s 104us/step - loss: 0.0141 - accuracy: 0.9954 - val_loss: 0.0736
- val_accuracy: 0.9815

Epoch 10/20

60000/60000 [=====] - 6s 105us/step - loss: 0.0117 - accuracy: 0.9959 - val_loss: 0.0878
- val_accuracy: 0.9797

Epoch 11/20

60000/60000 [=====] - 6s 106us/step - loss: 0.0138 - accuracy: 0.9957 - val_loss: 0.0906
- val_accuracy: 0.9788

Epoch 12/20

60000/60000 [=====] - 6s 105us/step - loss: 0.0115 - accuracy: 0.9960 - val_loss: 0.0676
- val_accuracy: 0.9839

Epoch 13/20

60000/60000 [=====] - 6s 103us/step - loss: 0.0104 - accuracy: 0.9963 - val_loss: 0.0970
- val_accuracy: 0.9782

Epoch 14/20

60000/60000 [=====] - 6s 104us/step - loss: 0.0099 - accuracy: 0.9965 - val_loss: 0.1023
- val_accuracy: 0.9792

Epoch 15/20

60000/60000 [=====] - 6s 105us/step - loss: 0.0111 - accuracy: 0.9964 - val_loss: 0.0926
- val_accuracy: 0.9803

Epoch 16/20

60000/60000 [=====] - 6s 105us/step - loss: 0.0071 - accuracy: 0.9976 - val_loss: 0.0873
- val_accuracy: 0.9805

Epoch 17/20

60000/60000 [=====] - 6s 104us/step - loss: 0.0078 - accuracy: 0.9972 - val_loss: 0.1013
- val_accuracy: 0.9788

Epoch 18/20

60000/60000 [=====] - 6s 104us/step - loss: 0.0096 - accuracy: 0.9968 - val_loss: 0.0973
- val_accuracy: 0.9822

Epoch 19/20

60000/60000 [=====] - 6s 105us/step - loss: 0.0063 - accuracy: 0.9979 - val_loss: 0.1016
- val_accuracy: 0.9811

Epoch 20/20

60000/60000 [=====] - 6s 103us/step - loss: 0.0049 - accuracy: 0.9985 - val_loss: 0.1079
- val_accuracy: 0.9796



```
In [33]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history1.history['val_loss']
ty = history1.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10785088968896661
Test accuracy: 0.9796000123023987



%matplotlib inline

plt.close('all')

1.2. MLP + ReLU + ADAM + Batch Normalization + HE normal initialization

```
In [36]: from keras.layers.normalization import BatchNormalization
        from keras.initializers import he_normal

        model_2 = Sequential()

        model_2.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
        model_2.add(Dense(128, activation='relu', kernel_initializer=he_normal(seed=None)) )
        model_2.add(BatchNormalization())

        model_2.add(Dense(output_dim, activation='softmax'))

        model_2.summary()

        model_2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

        history2 = model_2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test
, Y_test))
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 256)	200960
dense_12 (Dense)	(None, 128)	32896
batch_normalization_3 (Batch Normalization)	(None, 128)	512
dense_13 (Dense)	(None, 10)	1290

Total params: 235,658

Trainable params: 235,402

Non-trainable params: 256

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 58us/step - loss: 0.2228 - accuracy: 0.9351 - val_loss: 0.1234 - val_accuracy: 0.9608

Epoch 2/20

60000/60000 [=====] - 3s 54us/step - loss: 0.0859 - accuracy: 0.9741 - val_loss: 0.0993 - val_accuracy: 0.9699

Epoch 3/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0556 - accuracy: 0.9831 - val_loss: 0.0821 - val_accuracy: 0.9736

Epoch 4/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0405 - accuracy: 0.9874 - val_loss: 0.0836 - val_accuracy: 0.9733

Epoch 5/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0285 - accuracy: 0.9910 - val_loss: 0.0760 - val_accuracy: 0.9769

Epoch 6/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0219 - accuracy: 0.9929 - val_loss: 0.0801 - val_accuracy: 0.9769

Epoch 7/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0222 - accuracy: 0.9925 - val_loss: 0.0749 - val_accuracy: 0.9795

Epoch 8/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0179 - accuracy: 0.9940 - val_loss: 0.0730 - val_accuracy: 0.9796

Epoch 9/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0167 - accuracy: 0.9948 - val_loss: 0.0778 - val_accuracy: 0.9785

Epoch 10/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0092 - accuracy: 0.9971 - val_loss: 0.0807 - val_accuracy: 0.9790

Epoch 11/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0099 - accuracy: 0.9970 - val_loss: 0.0852 - val_accuracy: 0.9781

Epoch 12/20

60000/60000 [=====] - 3s 58us/step - loss: 0.0161 - accuracy: 0.9949 - val_loss: 0.0730 - val_accuracy: 0.9799

Epoch 13/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0113 - accuracy: 0.9962 - val_loss: 0.0860 - val_accuracy: 0.9796

Epoch 14/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0075 - accuracy: 0.9977 - val_loss: 0.0869 - val_accuracy: 0.9793

Epoch 15/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0070 - accuracy: 0.9977 - val_loss: 0.0905 - val_accuracy: 0.9783

Epoch 16/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0073 - accuracy: 0.9978 - val_loss: 0.0867 - val_accuracy: 0.9803

Epoch 17/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0091 - accuracy: 0.9969 - val_loss: 0.1147 - val_accuracy: 0.9765

Epoch 18/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0111 - accuracy: 0.9962 - val_loss: 0.0913 - val_accuracy: 0.9778

Epoch 19/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0049 - accuracy: 0.9985 - val_loss: 0.0907 - val_accuracy: 0.9791

Epoch 20/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0066 - accuracy: 0.9978 - val_loss: 0.1016 - val_accuracy: 0.9766


```
In [37]: score = model_2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history2.history['val_loss']
ty = history2.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10163859208421328
Test accuracy: 0.9765999913215637



1.3. MLP + ReLU + ADAM + Batch Normalization + Dropout + HE normal initialization

```
In [39]: from keras.layers.normalization import BatchNormalization
from keras.initializers import he_normal
from keras.layers import Dropout

model_3 = Sequential()

model_3.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
model_3.add(Dropout(0.5))

model_3.add(Dense(384, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_3.add(Dropout(0.5))

model_3.add(BatchNormalization())

model_3.add(Dense(output_dim, activation='softmax'))

model_3.summary()

model_3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history3 = model_3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test
, Y_test))
```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
dense_15 (Dense)	(None, 512)	401920
dropout_1 (Dropout)	(None, 512)	0
dense_16 (Dense)	(None, 384)	196992
dropout_2 (Dropout)	(None, 384)	0
batch_normalization_4 (Batch Normalization)	(None, 384)	1536
dense_17 (Dense)	(None, 10)	3850
Total params: 604,298		
Trainable params: 603,530		
Non-trainable params: 768		

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 9s 144us/step - loss: 0.4228 - accuracy: 0.8695 - val_loss: 0.1439
- val_accuracy: 0.9581

Epoch 2/20

60000/60000 [=====] - 8s 136us/step - loss: 0.1880 - accuracy: 0.9439 - val_loss: 0.0995
- val_accuracy: 0.9707

Epoch 3/20

60000/60000 [=====] - 8s 141us/step - loss: 0.1437 - accuracy: 0.9567 - val_loss: 0.0856
- val_accuracy: 0.9728

Epoch 4/20

60000/60000 [=====] - 8s 139us/step - loss: 0.1225 - accuracy: 0.9618 - val_loss: 0.0757
- val_accuracy: 0.9772

Epoch 5/20

60000/60000 [=====] - 8s 137us/step - loss: 0.1061 - accuracy: 0.9678 - val_loss: 0.0780
- val_accuracy: 0.9756

Epoch 6/20

60000/60000 [=====] - 8s 139us/step - loss: 0.0988 - accuracy: 0.9699 - val_loss: 0.0666
- val_accuracy: 0.9797

Epoch 7/20

60000/60000 [=====] - 8s 140us/step - loss: 0.0872 - accuracy: 0.9723 - val_loss: 0.0653
- val_accuracy: 0.9805

Epoch 8/20

60000/60000 [=====] - 8s 139us/step - loss: 0.0814 - accuracy: 0.9744 - val_loss: 0.0651
- val_accuracy: 0.9805

Epoch 9/20

60000/60000 [=====] - 8s 139us/step - loss: 0.0767 - accuracy: 0.9762 - val_loss: 0.0605
- val_accuracy: 0.9816

Epoch 10/20

60000/60000 [=====] - 8s 134us/step - loss: 0.0701 - accuracy: 0.9774 - val_loss: 0.0641
- val_accuracy: 0.9822

Epoch 11/20

60000/60000 [=====] - 8s 132us/step - loss: 0.0676 - accuracy: 0.9786 - val_loss: 0.0593
- val_accuracy: 0.9830

Epoch 12/20

60000/60000 [=====] - 8s 132us/step - loss: 0.0648 - accuracy: 0.9792 - val_loss: 0.0596
- val_accuracy: 0.9824

Epoch 13/20

60000/60000 [=====] - 8s 135us/step - loss: 0.0570 - accuracy: 0.9817 - val_loss: 0.0569
- val_accuracy: 0.9824

Epoch 14/20

60000/60000 [=====] - 8s 134us/step - loss: 0.0574 - accuracy: 0.9811 - val_loss: 0.0578
- val_accuracy: 0.9829

Epoch 15/20

60000/60000 [=====] - 8s 134us/step - loss: 0.0534 - accuracy: 0.9821 - val_loss: 0.0550
- val_accuracy: 0.9839

Epoch 16/20

60000/60000 [=====] - 8s 135us/step - loss: 0.0492 - accuracy: 0.9838 - val_loss: 0.0598
- val_accuracy: 0.9835

Epoch 17/20

60000/60000 [=====] - 8s 134us/step - loss: 0.0499 - accuracy: 0.9836 - val_loss: 0.0585
- val_accuracy: 0.9828

Epoch 18/20

60000/60000 [=====] - 8s 134us/step - loss: 0.0476 - accuracy: 0.9841 - val_loss: 0.0553
- val_accuracy: 0.9841

Epoch 19/20

60000/60000 [=====] - 8s 134us/step - loss: 0.0440 - accuracy: 0.9856 - val_loss: 0.0557
- val_accuracy: 0.9850

Epoch 20/20

60000/60000 [=====] - 8s 134us/step - loss: 0.0451 - accuracy: 0.9854 - val_loss: 0.0566
- val_accuracy: 0.9834

```
In [49]: score = model_3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history3.history['val_loss']
ty = history3.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.05660308131032507

Test accuracy: 0.9833999872207642



2. 3-Hidden layer architecture

2.1. MLP + ReLU + ADAM

```
In [41]: model_4 = Sequential()
model_4.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_4.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_4.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_4.add(Dense(output_dim, activation='softmax'))

print(model_4.summary())

model_4.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history4 = model_4.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```


Model: "sequential_9"

Layer (type)	Output Shape	Param #
dense_18 (Dense)	(None, 512)	401920
dense_19 (Dense)	(None, 256)	131328
dense_20 (Dense)	(None, 128)	32896
dense_21 (Dense)	(None, 10)	1290

Total params: 567,434
Trainable params: 567,434
Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 113us/step - loss: 0.2268 - accuracy: 0.9305 - val_loss: 0.1094
- val_accuracy: 0.9665

Epoch 2/20

60000/60000 [=====] - 7s 111us/step - loss: 0.0803 - accuracy: 0.9756 - val_loss: 0.0887
- val_accuracy: 0.9729

Epoch 3/20

60000/60000 [=====] - 7s 112us/step - loss: 0.0528 - accuracy: 0.9828 - val_loss: 0.0744
- val_accuracy: 0.9773

Epoch 4/20

60000/60000 [=====] - 7s 111us/step - loss: 0.0376 - accuracy: 0.9880 - val_loss: 0.0915
- val_accuracy: 0.9713

Epoch 5/20

60000/60000 [=====] - 7s 112us/step - loss: 0.0289 - accuracy: 0.9909 - val_loss: 0.0723
- val_accuracy: 0.9798

Epoch 6/20

60000/60000 [=====] - 7s 110us/step - loss: 0.0258 - accuracy: 0.9915 - val_loss: 0.0854
- val_accuracy: 0.9762

Epoch 7/20

60000/60000 [=====] - 7s 110us/step - loss: 0.0227 - accuracy: 0.9926 - val_loss: 0.0809
- val_accuracy: 0.9786

Epoch 8/20

60000/60000 [=====] - 7s 112us/step - loss: 0.0218 - accuracy: 0.9930 - val_loss: 0.0894
- val_accuracy: 0.9785

Epoch 9/20

60000/60000 [=====] - 7s 113us/step - loss: 0.0164 - accuracy: 0.9945 - val_loss: 0.0983
- val_accuracy: 0.9749

Epoch 10/20

60000/60000 [=====] - 7s 111us/step - loss: 0.0164 - accuracy: 0.9947 - val_loss: 0.0867
- val_accuracy: 0.9804

Epoch 11/20

60000/60000 [=====] - 6s 107us/step - loss: 0.0174 - accuracy: 0.9941 - val_loss: 0.0852
- val_accuracy: 0.9791

Epoch 12/20

60000/60000 [=====] - 6s 107us/step - loss: 0.0139 - accuracy: 0.9956 - val_loss: 0.0852
- val_accuracy: 0.9792

Epoch 13/20

60000/60000 [=====] - 6s 108us/step - loss: 0.0147 - accuracy: 0.9958 - val_loss: 0.1261
- val_accuracy: 0.9714

Epoch 14/20

60000/60000 [=====] - 7s 109us/step - loss: 0.0150 - accuracy: 0.9954 - val_loss: 0.0893
- val_accuracy: 0.9782

Epoch 15/20

60000/60000 [=====] - 7s 112us/step - loss: 0.0068 - accuracy: 0.9978 - val_loss: 0.0830
- val_accuracy: 0.9824

Epoch 16/20

60000/60000 [=====] - 7s 112us/step - loss: 0.0130 - accuracy: 0.9957 - val_loss: 0.0793
- val_accuracy: 0.9836

Epoch 17/20

60000/60000 [=====] - 6s 108us/step - loss: 0.0118 - accuracy: 0.9964 - val_loss: 0.0808
- val_accuracy: 0.9807

Epoch 18/20

60000/60000 [=====] - 7s 110us/step - loss: 0.0079 - accuracy: 0.9976 - val_loss: 0.0984
- val_accuracy: 0.9799

Epoch 19/20

60000/60000 [=====] - 7s 110us/step - loss: 0.0149 - accuracy: 0.9953 - val_loss: 0.1010
- val_accuracy: 0.9799

Epoch 20/20

60000/60000 [=====] - 7s 109us/step - loss: 0.0083 - accuracy: 0.9972 - val_loss: 0.1068
- val_accuracy: 0.9775


```
In [50]: score = model_4.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history4.history['val_loss']
ty = history4.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10676916894102763
Test accuracy: 0.9775000214576721



%matplotlib inline

plt.close('all')

1.2. MLP + ReLU + ADAM + Batch Normalization + HE normal initialization

```
In [44]: from keras.layers.normalization import BatchNormalization
        from keras.initializers import he_normal

        model_5 = Sequential()

        model_5.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
        model_5.add(Dense(384, activation='relu', kernel_initializer=he_normal(seed=None)) )
        model_5.add(Dense(512, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
        model_5.add(BatchNormalization())

        model_5.add(Dense(output_dim, activation='softmax'))

        model_5.summary()

        model_5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

        history5 = model_5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test
, Y_test))
```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
dense_22 (Dense)	(None, 256)	200960
dense_23 (Dense)	(None, 384)	98688
dense_24 (Dense)	(None, 512)	197120
batch_normalization_5 (Batch Normalization)	(None, 512)	2048
dense_25 (Dense)	(None, 10)	5130

Total params: 503,946

Trainable params: 502,922

Non-trainable params: 1,024

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 121us/step - loss: 0.2097 - accuracy: 0.9366 - val_loss: 0.1158
- val_accuracy: 0.9610

Epoch 2/20

60000/60000 [=====] - 7s 120us/step - loss: 0.0848 - accuracy: 0.9737 - val_loss: 0.1262
- val_accuracy: 0.9619

Epoch 3/20

60000/60000 [=====] - 7s 120us/step - loss: 0.0627 - accuracy: 0.9805 - val_loss: 0.0781
- val_accuracy: 0.9760

Epoch 4/20

60000/60000 [=====] - 7s 120us/step - loss: 0.0477 - accuracy: 0.9846 - val_loss: 0.0862
- val_accuracy: 0.9738

Epoch 5/20

60000/60000 [=====] - 7s 121us/step - loss: 0.0351 - accuracy: 0.9883 - val_loss: 0.1040
- val_accuracy: 0.9689

Epoch 6/20

60000/60000 [=====] - 7s 118us/step - loss: 0.0316 - accuracy: 0.9893 - val_loss: 0.0931
- val_accuracy: 0.9733

Epoch 7/20

60000/60000 [=====] - 7s 118us/step - loss: 0.0261 - accuracy: 0.9916 - val_loss: 0.0805
- val_accuracy: 0.9796

Epoch 8/20

60000/60000 [=====] - 7s 119us/step - loss: 0.0210 - accuracy: 0.9929 - val_loss: 0.0913
- val_accuracy: 0.9755

Epoch 9/20

60000/60000 [=====] - 7s 121us/step - loss: 0.0207 - accuracy: 0.9930 - val_loss: 0.1076
- val_accuracy: 0.9734

Epoch 10/20

60000/60000 [=====] - 7s 125us/step - loss: 0.0182 - accuracy: 0.9939 - val_loss: 0.0949
- val_accuracy: 0.9745

Epoch 11/20

60000/60000 [=====] - 8s 127us/step - loss: 0.0157 - accuracy: 0.9945 - val_loss: 0.0788
- val_accuracy: 0.9782

Epoch 12/20

60000/60000 [=====] - 8s 129us/step - loss: 0.0167 - accuracy: 0.9945 - val_loss: 0.0946
- val_accuracy: 0.9779

Epoch 13/20

60000/60000 [=====] - 8s 131us/step - loss: 0.0181 - accuracy: 0.9943 - val_loss: 0.0817
- val_accuracy: 0.9799

Epoch 14/20

60000/60000 [=====] - 7s 121us/step - loss: 0.0115 - accuracy: 0.9961 - val_loss: 0.0830
- val_accuracy: 0.9794

Epoch 15/20

60000/60000 [=====] - 7s 122us/step - loss: 0.0115 - accuracy: 0.9963 - val_loss: 0.0980
- val_accuracy: 0.9766

Epoch 16/20

60000/60000 [=====] - 7s 123us/step - loss: 0.0113 - accuracy: 0.9963 - val_loss: 0.1014
- val_accuracy: 0.9783

Epoch 17/20

60000/60000 [=====] - 7s 123us/step - loss: 0.0127 - accuracy: 0.9960 - val_loss: 0.1022
- val_accuracy: 0.9784

Epoch 18/20

60000/60000 [=====] - 7s 122us/step - loss: 0.0105 - accuracy: 0.9965 - val_loss: 0.0914
- val_accuracy: 0.9793

Epoch 19/20

60000/60000 [=====] - 7s 121us/step - loss: 0.0094 - accuracy: 0.9966 - val_loss: 0.1374
- val_accuracy: 0.9735

Epoch 20/20

60000/60000 [=====] - 7s 122us/step - loss: 0.0124 - accuracy: 0.9959 - val_loss: 0.1005
- val_accuracy: 0.9800


```
In [51]: score = model_5.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history5.history['val_loss']
ty = history5.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10054782296594231
Test accuracy: 0.980000190734863



1.3. MLP + ReLU + ADAM + Batch Normalization + Dropout + HE normal initialization

```
In [46]: from keras.layers.normalization import BatchNormalization
from keras.initializers import he_normal
from keras.layers import Dropout

model_6 = Sequential()

model_6.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
model_6.add(BatchNormalization())
model_6.add(Dropout(0.3))

model_6.add(Dense(384, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_6.add(BatchNormalization())
model_6.add(Dropout(0.4))

model_6.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_6.add(BatchNormalization())
model_6.add(Dropout(0.5))

model_6.add(Dense(output_dim, activation='softmax'))

model_6.summary()

model_6.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history6 = model_6.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test
, Y_test))
```

Model: "sequential_11"

Layer (type)	Output Shape	Param #
dense_26 (Dense)	(None, 256)	200960
batch_normalization_6 (Batch Normalization)	(None, 256)	1024
dropout_3 (Dropout)	(None, 256)	0
dense_27 (Dense)	(None, 384)	98688
batch_normalization_7 (Batch Normalization)	(None, 384)	1536
dropout_4 (Dropout)	(None, 384)	0
dense_28 (Dense)	(None, 256)	98560
batch_normalization_8 (Batch Normalization)	(None, 256)	1024
dropout_5 (Dropout)	(None, 256)	0
dense_29 (Dense)	(None, 10)	2570

=====
Total params: 404,362
Trainable params: 402,570
Non-trainable params: 1,792

=====
Train on 60000 samples, validate on 10000 samples

Epoch 1/20
60000/60000 [=====] - 8s 125us/step - loss: 0.4752 - accuracy: 0.8572 - val_loss: 0.1460
- val_accuracy: 0.9537
Epoch 2/20
60000/60000 [=====] - 7s 115us/step - loss: 0.2098 - accuracy: 0.9369 - val_loss: 0.1075
- val_accuracy: 0.9673
Epoch 3/20
60000/60000 [=====] - 7s 116us/step - loss: 0.1569 - accuracy: 0.9522 - val_loss: 0.0933
- val_accuracy: 0.9719
Epoch 4/20
60000/60000 [=====] - 7s 117us/step - loss: 0.1319 - accuracy: 0.9596 - val_loss: 0.0828
- val_accuracy: 0.9727
Epoch 5/20
60000/60000 [=====] - 7s 116us/step - loss: 0.1123 - accuracy: 0.9656 - val_loss: 0.0785
- val_accuracy: 0.9764
Epoch 6/20
60000/60000 [=====] - 7s 117us/step - loss: 0.1039 - accuracy: 0.9676 - val_loss: 0.0812
- val_accuracy: 0.9757
Epoch 7/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0946 - accuracy: 0.9713 - val_loss: 0.0721
- val_accuracy: 0.9773
Epoch 8/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0840 - accuracy: 0.9739 - val_loss: 0.0687
- val_accuracy: 0.9787
Epoch 9/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0800 - accuracy: 0.9747 - val_loss: 0.0696
- val_accuracy: 0.9788
Epoch 10/20
60000/60000 [=====] - 7s 117us/step - loss: 0.0735 - accuracy: 0.9770 - val_loss: 0.0661
- val_accuracy: 0.9815
Epoch 11/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0703 - accuracy: 0.9779 - val_loss: 0.0590
- val_accuracy: 0.9814
Epoch 12/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0661 - accuracy: 0.9791 - val_loss: 0.0628
- val_accuracy: 0.9803
Epoch 13/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0622 - accuracy: 0.9804 - val_loss: 0.0603
- val_accuracy: 0.9811
Epoch 14/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0582 - accuracy: 0.9822 - val_loss: 0.0594
- val_accuracy: 0.9835
Epoch 15/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0544 - accuracy: 0.9828 - val_loss: 0.0548
- val_accuracy: 0.9849
Epoch 16/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0528 - accuracy: 0.9826 - val_loss: 0.0594
- val_accuracy: 0.9825
Epoch 17/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0506 - accuracy: 0.9834 - val_loss: 0.0605
- val_accuracy: 0.9827
Epoch 18/20
60000/60000 [=====] - 7s 117us/step - loss: 0.0486 - accuracy: 0.9838 - val_loss: 0.0573
- val_accuracy: 0.9827
Epoch 19/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0480 - accuracy: 0.9852 - val_loss: 0.0561
- val_accuracy: 0.9843
Epoch 20/20

60000/60000 [=====] - 7s 117us/step - loss: 0.0415 - accuracy: 0.9870 - val_loss: 0.0638
- val_accuracy: 0.9829


```
In [52]: score = model_6.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history6.history['val_loss']
ty = history6.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06382637676765153
Test accuracy: 0.9829000234603882



In [0]:

3. 5-Hidden layer architecture

3.1. MLP + ReLU + ADAM


```
In [48]: model_7 = Sequential()
model_7.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_7.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_7.add(Dense(512, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_7.add(Dense(256, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_7.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_7.add(Dense(output_dim, activation='softmax'))

print(model_7.summary())

model_7.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history7 = model_7.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Model: "sequential_12"

Layer (type)	Output Shape	Param #
dense_30 (Dense)	(None, 128)	100480
dense_31 (Dense)	(None, 256)	33024
dense_32 (Dense)	(None, 512)	131584
dense_33 (Dense)	(None, 256)	131328
dense_34 (Dense)	(None, 128)	32896
dense_35 (Dense)	(None, 10)	1290

=====
Total params: 430,602
Trainable params: 430,602
Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 101us/step - loss: 0.2822 - accuracy: 0.9142 - val_loss: 0.1245
- val_accuracy: 0.9630

Epoch 2/20

60000/60000 [=====] - 6s 100us/step - loss: 0.1103 - accuracy: 0.9662 - val_loss: 0.1154
- val_accuracy: 0.9662

Epoch 3/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0743 - accuracy: 0.9771 - val_loss: 0.0957 -
val_accuracy: 0.9721

Epoch 4/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0599 - accuracy: 0.9810 - val_loss: 0.0966 -
val_accuracy: 0.9720

Epoch 5/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0456 - accuracy: 0.9854 - val_loss: 0.1049 -
val_accuracy: 0.9707

Epoch 6/20

60000/60000 [=====] - 6s 101us/step - loss: 0.0402 - accuracy: 0.9873 - val_loss: 0.0943
- val_accuracy: 0.9723

Epoch 7/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0369 - accuracy: 0.9880 - val_loss: 0.0884 -
val_accuracy: 0.9753

Epoch 8/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0333 - accuracy: 0.9895 - val_loss: 0.1322
- val_accuracy: 0.9684

Epoch 9/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0282 - accuracy: 0.9912 - val_loss: 0.1010 -
val_accuracy: 0.9741

Epoch 10/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0281 - accuracy: 0.9911 - val_loss: 0.1097 -
val_accuracy: 0.9757

Epoch 11/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0235 - accuracy: 0.9927 - val_loss: 0.1033 -
val_accuracy: 0.9743

Epoch 12/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0234 - accuracy: 0.9927 - val_loss: 0.1020 -
val_accuracy: 0.9764

Epoch 13/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0199 - accuracy: 0.9935 - val_loss: 0.1223 -
val_accuracy: 0.9738

Epoch 14/20

60000/60000 [=====] - 6s 101us/step - loss: 0.0219 - accuracy: 0.9931 - val_loss: 0.1211
- val_accuracy: 0.9751

Epoch 15/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0169 - accuracy: 0.9946 - val_loss: 0.1082
- val_accuracy: 0.9751

Epoch 16/20

60000/60000 [=====] - 6s 97us/step - loss: 0.0190 - accuracy: 0.9941 - val_loss: 0.1063 -
val_accuracy: 0.9767

Epoch 17/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0157 - accuracy: 0.9952 - val_loss: 0.1223 -
val_accuracy: 0.9736

Epoch 18/20

60000/60000 [=====] - 6s 105us/step - loss: 0.0171 - accuracy: 0.9947 - val_loss: 0.1059
- val_accuracy: 0.9779

Epoch 19/20

60000/60000 [=====] - 6s 103us/step - loss: 0.0150 - accuracy: 0.9956 - val_loss: 0.1050
- val_accuracy: 0.9779

Epoch 20/20

60000/60000 [=====] - 6s 97us/step - loss: 0.0178 - accuracy: 0.9948 - val_loss: 0.1090 -
val_accuracy: 0.9771


```
In [53]: score = model_7.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history7.history['val_loss']
ty = history7.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10898555530682802
Test accuracy: 0.9771000146865845



%matplotlib inline

plt.close('all')

3.2. MLP + ReLU + ADAM + Batch Normalization + HE normal initialization

```
In [54]: from keras.layers.normalization import BatchNormalization
        from keras.initializers import he_normal

        model_8 = Sequential()

        model_8.add(Dense(128, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
        model_8.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None)) )
        model_8.add(BatchNormalization())
        model_8.add(Dense(384, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
        model_8.add(BatchNormalization())
        model_8.add(Dense(512, activation='relu', kernel_initializer=he_normal(seed=None)) )
        model_8.add(BatchNormalization())
        model_8.add(Dense(384, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
        model_8.add(BatchNormalization())

        model_8.add(Dense(output_dim, activation='softmax'))

        model_8.summary()

        model_8.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

        history8 = model_8.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test
, Y_test))
```

Model: "sequential_13"

Layer (type)	Output Shape	Param #
dense_36 (Dense)	(None, 128)	100480
dense_37 (Dense)	(None, 256)	33024
batch_normalization_9 (Batch Normalization)	(None, 256)	1024
dense_38 (Dense)	(None, 384)	98688
batch_normalization_10 (Batch Normalization)	(None, 384)	1536
dense_39 (Dense)	(None, 512)	197120
batch_normalization_11 (Batch Normalization)	(None, 512)	2048
dense_40 (Dense)	(None, 384)	196992
batch_normalization_12 (Batch Normalization)	(None, 384)	1536
dense_41 (Dense)	(None, 10)	3850

=====
Total params: 636,298
Trainable params: 633,226
Non-trainable params: 3,072

=====
Train on 60000 samples, validate on 10000 samples

Epoch 1/20
60000/60000 [=====] - 10s 163us/step - loss: 0.2359 - accuracy: 0.9273 - val_loss: 0.1322
- val_accuracy: 0.9598
Epoch 2/20
60000/60000 [=====] - 9s 157us/step - loss: 0.0929 - accuracy: 0.9713 - val_loss: 0.1262
- val_accuracy: 0.9613
Epoch 3/20
60000/60000 [=====] - 9s 155us/step - loss: 0.0638 - accuracy: 0.9793 - val_loss: 0.1127
- val_accuracy: 0.9660
Epoch 4/20
60000/60000 [=====] - 9s 150us/step - loss: 0.0509 - accuracy: 0.9826 - val_loss: 0.1158
- val_accuracy: 0.9670
Epoch 5/20
60000/60000 [=====] - 9s 148us/step - loss: 0.0412 - accuracy: 0.9856 - val_loss: 0.1240
- val_accuracy: 0.9640
Epoch 6/20
60000/60000 [=====] - 9s 150us/step - loss: 0.0378 - accuracy: 0.9873 - val_loss: 0.0794
- val_accuracy: 0.9740
Epoch 7/20
60000/60000 [=====] - 9s 151us/step - loss: 0.0318 - accuracy: 0.9888 - val_loss: 0.0763
- val_accuracy: 0.9777
Epoch 8/20
60000/60000 [=====] - 9s 150us/step - loss: 0.0317 - accuracy: 0.9890 - val_loss: 0.0817
- val_accuracy: 0.9769
Epoch 9/20
60000/60000 [=====] - 9s 149us/step - loss: 0.0282 - accuracy: 0.9908 - val_loss: 0.0792
- val_accuracy: 0.9782
Epoch 10/20
60000/60000 [=====] - 9s 149us/step - loss: 0.0256 - accuracy: 0.9912 - val_loss: 0.0851
- val_accuracy: 0.9773
Epoch 11/20
60000/60000 [=====] - 9s 146us/step - loss: 0.0232 - accuracy: 0.9921 - val_loss: 0.0930
- val_accuracy: 0.9760
Epoch 12/20
60000/60000 [=====] - 9s 149us/step - loss: 0.0199 - accuracy: 0.9931 - val_loss: 0.0995
- val_accuracy: 0.9756
Epoch 13/20
60000/60000 [=====] - 9s 152us/step - loss: 0.0215 - accuracy: 0.9925 - val_loss: 0.0923
- val_accuracy: 0.9756
Epoch 14/20
60000/60000 [=====] - 9s 152us/step - loss: 0.0216 - accuracy: 0.9929 - val_loss: 0.1121
- val_accuracy: 0.9732
Epoch 15/20
60000/60000 [=====] - 9s 149us/step - loss: 0.0168 - accuracy: 0.9943 - val_loss: 0.0971
- val_accuracy: 0.9753
Epoch 16/20
60000/60000 [=====] - 9s 156us/step - loss: 0.0163 - accuracy: 0.9945 - val_loss: 0.1065
- val_accuracy: 0.9747
Epoch 17/20
60000/60000 [=====] - 10s 160us/step - loss: 0.0195 - accuracy: 0.9937 - val_loss: 0.0850
- val_accuracy: 0.9797
Epoch 18/20
60000/60000 [=====] - 10s 163us/step - loss: 0.0147 - accuracy: 0.9957 - val_loss: 0.0887
- val_accuracy: 0.9788
Epoch 19/20
60000/60000 [=====] - 10s 161us/step - loss: 0.0152 - accuracy: 0.9952 - val_loss: 0.0815
- val_accuracy: 0.9797
Epoch 20/20

60000/60000 [=====] - 10s 160us/step - loss: 0.0126 - accuracy: 0.9959 - val_loss: 0.0885
- val_accuracy: 0.9790


```
In [55]: score = model_8.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history8.history['val_loss']
ty = history8.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08854798433639116
Test accuracy: 0.9789999723434448



3.3. MLP + ReLU + ADAM + Batch Normalization + Dropout + HE normal initialization

```
In [57]: from keras.layers.normalization import BatchNormalization
from keras.initializers import he_normal
from keras.layers import Dropout

model_9 = Sequential()

model_9.add(Dense(256, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
model_9.add(BatchNormalization())
model_9.add(Dropout(0.3))

model_9.add(Dense(384, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_9.add(BatchNormalization())
model_9.add(Dropout(0.3))

model_9.add(Dense(512, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_9.add(BatchNormalization())
model_9.add(Dropout(0.3))

model_9.add(Dense(384, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_9.add(BatchNormalization())
model_9.add(Dropout(0.3))

model_9.add(Dense(256, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_9.add(BatchNormalization())
model_9.add(Dropout(0.3))

model_9.add(Dense(output_dim, activation='softmax'))

model_9.summary()

model_9.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history9 = model_9.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Model: "sequential_15"

Layer (type)	Output Shape	Param #
dense_46 (Dense)	(None, 256)	200960
batch_normalization_16 (Batch Normalization)	(None, 256)	1024
dropout_9 (Dropout)	(None, 256)	0
dense_47 (Dense)	(None, 384)	98688
batch_normalization_17 (Batch Normalization)	(None, 384)	1536
dropout_10 (Dropout)	(None, 384)	0
dense_48 (Dense)	(None, 512)	197120
batch_normalization_18 (Batch Normalization)	(None, 512)	2048
dropout_11 (Dropout)	(None, 512)	0
dense_49 (Dense)	(None, 384)	196992
batch_normalization_19 (Batch Normalization)	(None, 384)	1536
dropout_12 (Dropout)	(None, 384)	0
dense_50 (Dense)	(None, 256)	98560
batch_normalization_20 (Batch Normalization)	(None, 256)	1024
dropout_13 (Dropout)	(None, 256)	0
dense_51 (Dense)	(None, 10)	2570

=====
Total params: 802,058
Trainable params: 798,474
Non-trainable params: 3,584

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 13s 218us/step - loss: 0.4796 - accuracy: 0.8537 - val_loss: 0.1411
- val_accuracy: 0.9555

Epoch 2/20

60000/60000 [=====] - 12s 204us/step - loss: 0.2065 - accuracy: 0.9383 - val_loss: 0.1054
- val_accuracy: 0.9679

Epoch 3/20

60000/60000 [=====] - 12s 207us/step - loss: 0.1593 - accuracy: 0.9514 - val_loss: 0.0895
- val_accuracy: 0.9733

Epoch 4/20

60000/60000 [=====] - 12s 201us/step - loss: 0.1370 - accuracy: 0.9571 - val_loss: 0.0891
- val_accuracy: 0.9733

Epoch 5/20

60000/60000 [=====] - 13s 212us/step - loss: 0.1178 - accuracy: 0.9639 - val_loss: 0.0760
- val_accuracy: 0.9773

Epoch 6/20

60000/60000 [=====] - 12s 203us/step - loss: 0.1082 - accuracy: 0.9677 - val_loss: 0.0766
- val_accuracy: 0.9766

Epoch 7/20

60000/60000 [=====] - 12s 200us/step - loss: 0.0964 - accuracy: 0.9698 - val_loss: 0.0726
- val_accuracy: 0.9776

Epoch 8/20

60000/60000 [=====] - 12s 199us/step - loss: 0.0913 - accuracy: 0.9722 - val_loss: 0.0757
- val_accuracy: 0.9767

Epoch 9/20

60000/60000 [=====] - 12s 199us/step - loss: 0.0860 - accuracy: 0.9728 - val_loss: 0.0779
- val_accuracy: 0.9776

Epoch 10/20

60000/60000 [=====] - 12s 201us/step - loss: 0.0809 - accuracy: 0.9753 - val_loss: 0.0684
- val_accuracy: 0.9796

Epoch 11/20

60000/60000 [=====] - 12s 195us/step - loss: 0.0759 - accuracy: 0.9755 - val_loss: 0.0673
- val_accuracy: 0.9796

Epoch 12/20

60000/60000 [=====] - 12s 194us/step - loss: 0.0718 - accuracy: 0.9772 - val_loss: 0.0732
- val_accuracy: 0.9776

Epoch 13/20

60000/60000 [=====] - 12s 194us/step - loss: 0.0687 - accuracy: 0.9784 - val_loss: 0.0665
- val_accuracy: 0.9786

Epoch 14/20

60000/60000 [=====] - 12s 200us/step - loss: 0.0611 - accuracy: 0.9804 - val_loss: 0.0631
- val_accuracy: 0.9820

Epoch 15/20

60000/60000 [=====] - 12s 199us/step - loss: 0.0602 - accuracy: 0.9814 - val_loss: 0.0641
- val_accuracy: 0.9819

Epoch 16/20


```

60000/60000 [=====] - 12s 196us/step - loss: 0.0574 - accuracy: 0.9822 - val_loss: 0.0645
- val_accuracy: 0.9820
Epoch 17/20
60000/60000 [=====] - 12s 204us/step - loss: 0.0550 - accuracy: 0.9827 - val_loss: 0.0563
- val_accuracy: 0.9837
Epoch 18/20
60000/60000 [=====] - 13s 214us/step - loss: 0.0539 - accuracy: 0.9829 - val_loss: 0.0619
- val_accuracy: 0.9841
Epoch 19/20
60000/60000 [=====] - 13s 215us/step - loss: 0.0513 - accuracy: 0.9838 - val_loss: 0.0584
- val_accuracy: 0.9840
Epoch 20/20
60000/60000 [=====] - 13s 218us/step - loss: 0.0516 - accuracy: 0.9835 - val_loss: 0.0643
- val_accuracy: 0.9817

```

```

In [58]: score = model_9.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# List of epoch numbers
x = list(range(1,nb_epoch+1))

vy = history9.history['val_loss']
ty = history9.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06429058624608443
Test accuracy: 0.9817000031471252



```

In [64]: from prettytable import PrettyTable
x = PrettyTable()
x.field_names = ["Hidden Layers","Parameters","Loss"]
x.add_row([2,'MLP + ReLU + ADAM ', 0.979])
x.add_row([2,'MLP + ReLU + ADAM + Batch Normalization + HE normal initialization', 0.976])
x.add_row([2,'MLP + ReLU + ADAM + Batch Normalization + Dropout + HE normal initialization', 0.983])
x.add_row([3,'MLP + ReLU + ADAM ', 0.977])
x.add_row([3,'MLP + ReLU + ADAM + Batch Normalization + HE normal initialization', 0.980])
x.add_row([3,'MLP + ReLU + ADAM + Batch Normalization + Dropout + HE normal initialization', 0.982])
x.add_row([5,'MLP + ReLU + ADAM ', 0.977])
x.add_row([5,'MLP + ReLU + ADAM + Batch Normalization + HE normal initialization', 0.978])
x.add_row([5,'MLP + ReLU + ADAM + Batch Normalization + Dropout + HE normal initialization', 0.981])
print(x)

```

Hidden Layers	Parameters	Loss
2	MLP + ReLU + ADAM	0.979
2	MLP + ReLU + ADAM + Batch Normalization + HE normal initialization	0.976
2	MLP + ReLU + ADAM + Batch Normalization + Dropout + HE normal initialization	0.983
3	MLP + ReLU + ADAM	0.977
3	MLP + ReLU + ADAM + Batch Normalization + HE normal initialization	0.980
3	MLP + ReLU + ADAM + Batch Normalization + Dropout + HE normal initialization	0.982
5	MLP + ReLU + ADAM	0.977
5	MLP + ReLU + ADAM + Batch Normalization + HE normal initialization	0.978
5	MLP + ReLU + ADAM + Batch Normalization + Dropout + HE normal initialization	0.981

In [0]: